



AUTOMATING TREND DISCOVERY

On Streaming Datasets with Apache Spark

#streamingdiscovery | @newfront | @odsc

Introductions

Scott Haines

[@newfront](#)

newfrontcreative@gmail.com



<https://github.com/newfront/odsc-west-streaming-trends>

Workshop Goals

Data Architecture / Fall
back in love with your
data



Workshop Goals

Use Spark to Clean and
Explore Data



Workshop Goals

Learn how to harness
Trend Discovery / Why
unsupervised matters

Walk through a real
application



Workshop Goals

Understand how to test
Spark Applications and
how this makes
shipping to prod great!



Workshop Goals

Simplifying Monitoring




PART ONE: DATA ARCHITECTURE

What is Data Architecture?

#streamingdiscovery | @newfront | @odsc

Data Architecture


Data Lake Style



Makes Sense in Theory

Data Architecture

Data Lake Style




select * from table where **a IS NOT NULL** and **b IS NOT NULL...**

Extracting the data is a
different story all together

Data Architecture

Data Lake Style

1. Data is Volatile.
2. Field Types can change.
3. Older Data can become broken...



In Reality...

PART ONE: DATA ARCHITECTURE

Establish a Data Contract

#streamingdiscovery | @newfront | @odsc

Data Architecture


Data Contract

1. Define Data Rules
2. Define Types
3. Use Interoperable
Data Formats
4. Validate Data
Completeness

```
message UserEvent {  
  required uint32 schema_version = 1;  
  required bool gdpr_redacted      = 2;  
  required string user_id          = 3;  
  required string uuid             = 4;  
  required uint64 event_ts         = 5;  
  uint64 logged_event_ts           = 6;  
  required UserEventType event     = 7;  
  required UserEventData data      = 8;  
}
```

Know what to expect

PART ONE: DATA ARCHITECTURE PIPELINE



/api/v1/user/events


/api/v2/user/events

#streamingdiscovery | @newfront | @odsc

Data Architecture

Structured Data Store

1. Clean Valid Data
2. Authenticated before ingestion
3. Backwards Compatible



Valid, Clean Data


Data Architecture

Structured Data Store


1. Can Sanitize and Store in GDPR compliant Democratized All-Access store too



V2



Redaction Filter



Valid, Compliant, Clean Data

FALL BACK IN LOVE WITH YOUR DATA ARCHITECTURE

Questions?

#streamingdiscovery | @newfront | @odsc

First Steps

Quick Spark 101

1. DataFrames / Datasets

DataFrame's have Structure!

```
val us = Seq(  
  ("Scott", "Teacher"),  
  ("You", "Student")  
)  
.toDF("name", "role")  
  
us.printSchema  
/*  
root  
|-- name: string (nullable = true)  
|-- role: string (nullable = true)  
*/
```

First Steps

DataFrames/Datasets

Easily joined, mutated and aggregated

```
coffeeStand
```

```
.join(coffeeRatings, coffeeStand("name") === coffeeRatings("coffeeName"))  
.drop("coffeeName")  
.groupBy("name")  
.agg(avg("score") as "rating")  
.sort(desc("rating"))
```

First Steps

DataFrames/Datasets

Did I mention that all of this can be done streaming? Let's take a peek.

/part2/coffee/streaming_coffee.scala

Your First Streaming App

```
import spark.implicits._
```

```
def asCoffeeRating(input: String): CoffeeRating = {  
  val data = input.split(",")  
  val coffeeName = data(0)  
  val score = data(1).toInt  
  val note = if (data.size > 2) Some(data(2)) else None  
  CoffeeRating(coffeeName, score, note)  
}
```

```
val coffeeStandDF = sparkSession.sparkContext.parallelize(availableCoffee, 3).toDF  
val coffeeRatingsReader = sparkSession.readStream.format("socket").option("host", "localhost").option("port", 9999).load()  
val rawRatingsData: Dataset[String] = coffeeRatingsReader.as[String]  
  
val coffeeRatingsInput = rawRatingsData.map { asCoffeeRating }.toDF  
val coffeeAndRatingsDF = coffeeStandDF.join(coffeeRatingsInput, coffeeStandDF("name") === coffeeRatingsInput("coffeeName"))  
val averageRatings = coffeeAndRatingsDF.groupBy(col("name")).agg(avg("score") as "rating").sort(desc("rating"))  
val query = averageRatings.writeStream.outputMode("complete").format("console").start()
```

/part2/coffee/streaming_coffee.scala

First Steps

Code Walk Through

/part2/coffee/

FIRST STEPS: DOING SOME THINGS WITH SPARK

Questions?

#streamingdiscovery | @newfront | @odsc

PART TWO: CLEAN AND EXPLORE YOUR DATA


Cause we don't all have Data Contracts!

#streamingdiscovery | @newfront | @odsc

Data Analysis

Core Concepts

1. Loading
2. Exploring
3. Cleaning
4. Filling
5. More Exploring
6. Apriori
7. KMeans



Let's Play with Wine Reviews

/part2/wine/hello-wine.scala

/part2/wine/wine_reviews.scala

/part2/wine/wine_reviews_json.scala

Data Cleaning and Analysis

Code Walk Through

/part2/wine/

PART TWO: CLEAN AND EXPLORE YOUR DATA

Questions?

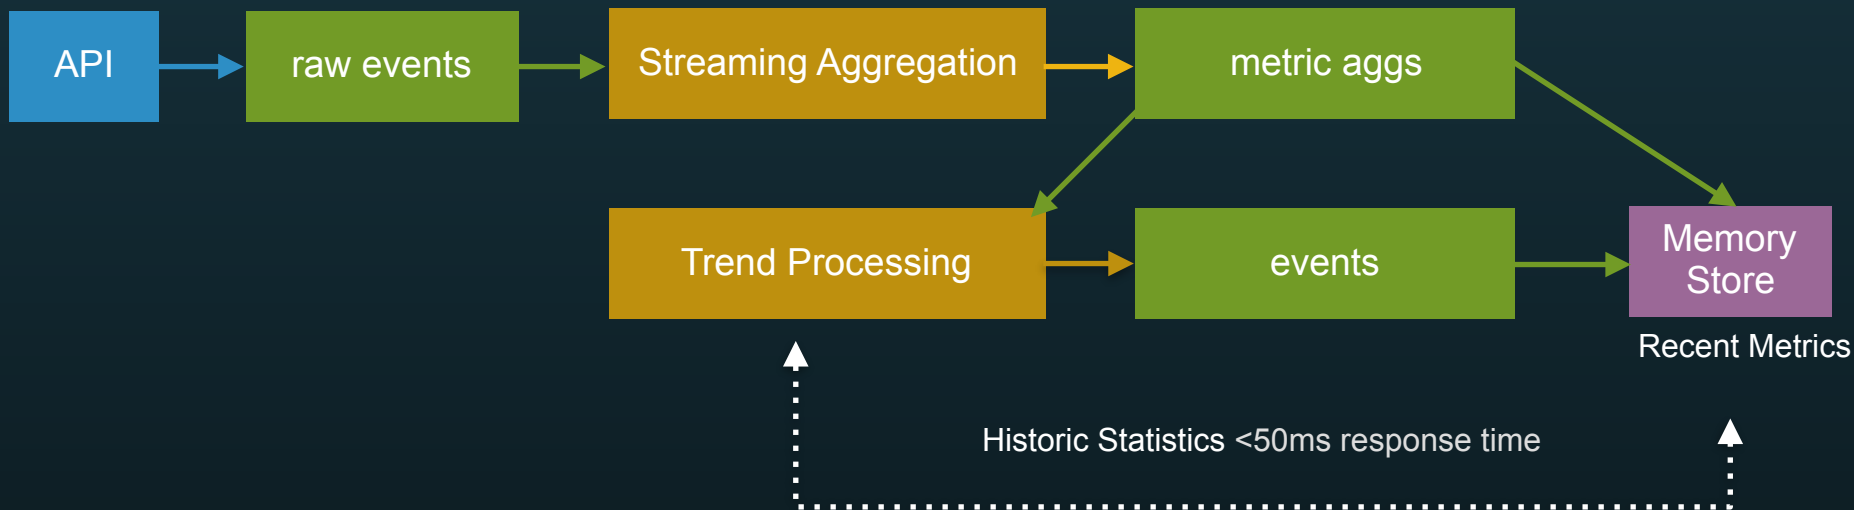
#streamingdiscovery | @newfront | @odsc

PART THREE: TREND DISCOVERY

Using statistics to understand unlabeled
trends in streaming datasets

#streamingdiscovery | @newfront | @odsc

STREAMING TREND DISCOVERY: ARCHITECTURE




Trend Discovery

Core Concepts

1. Kafka Streaming
2. Windowing
3. Watermarking
4. FlatMapGroupsWithState
5. DataSketching
6. Monadic Systems
7. And More!

Let's graduate from the Shell...



/part3/streaming-trend-discovery

Trend Discovery

Data Model

1. Generic Structures are nice.

```
message MetricAggregation {  
    optional string metric                = 1;  
    optional uint64 window_start          = 2;  
    optional uint64 window_end            = 3;  
    optional string window_interval       = 4;  
    optional uint32 samples                = 5;  
    optional Stats stats                  = 6;  
    optional Histogram histogram          = 7;  
    optional Dimensions dimensions        = 8;  
    optional string dimension_hash         = 10;  
}
```

/part3/streaming-trend-discovery

Trend Discovery

Data Model

1. Standardize on conventions

```
message Window {  
  optional string start      = 1;  
  optional uint64 start_ms   = 2;  
  optional string interval   = 3;  
  optional string end        = 4;  
  optional uint64 end_ms     = 5;  
}
```

/part3/streaming-trend-discovery

Trend Discovery

Data Model

1. Solve common problems

```
message Stats {  
  optional double min      = 1;  
  optional double p25      = 2;  
  optional double median   = 3;  
  optional double p75      = 4;  
  optional double p90      = 5;  
  optional double p95      = 6;  
  optional double p99      = 7;  
  optional double max      = 8;  
  optional double mean     = 9;  
  optional double sd       = 10;  
  optional double variance = 11;  
}
```

/part3/streaming-trend-discovery

Trend Discovery

Data Model

1. Names should describe the data for humans.

```
message Metric {  
  optional uint64 timestamp      = 1;  
  optional string group_key      = 2;  
  optional string dimensional_hash = 3;  
  optional string metric_name    = 4;  
  optional string label          = 5;  
  optional float value           = 6;  
  optional string carrier        = 7;  
  optional string country        = 8;  
  optional string route          = 9;  
  optional string direction      = 10;  
}
```

/part3/streaming-trend-discovery

Trend Discovery

Windowing / Watermark

1. Names should describe the data for humans.

```
callRecords
```

```
.withWatermark("timestamp", s"$watermarkInterval")  
.dropDuplicates("callSid", "timestamp")  
.groupBy($"accountSid", window($"timestamp", s"$windowInterval"))  
.agg(  
  min("pdd") as "minPdd",  
  round(avg("pdd"), 2) as "avgPdd",  
  max("pdd") as "maxPdd",  
  min("duration") as "minDuration",  
  round(avg("duration"), 2) as "avgDuration",  
  max("duration") as "maxDuration",  
  count("callSid") as "calls",  
  CallStateAggregation($"callState") as "callState",  
  DirectionAggregation($"direction") as "direction",  
  DisconnectedByAggregation($"disconnectedBy") as "disconnectedBy",  
  LastSipResponseAggregation($"lastSipResponse") as "lastSipResponse",  
  ProvidersAggregation($"providerSid") as "providerSid",  
  CountriesAggregation($"callerCountry") as "callerCountry",  
  CountriesAggregation($"calleeCountry") as "calleeCountry"  
)
```

/part3/streaming-trend-discovery

Trend Discovery

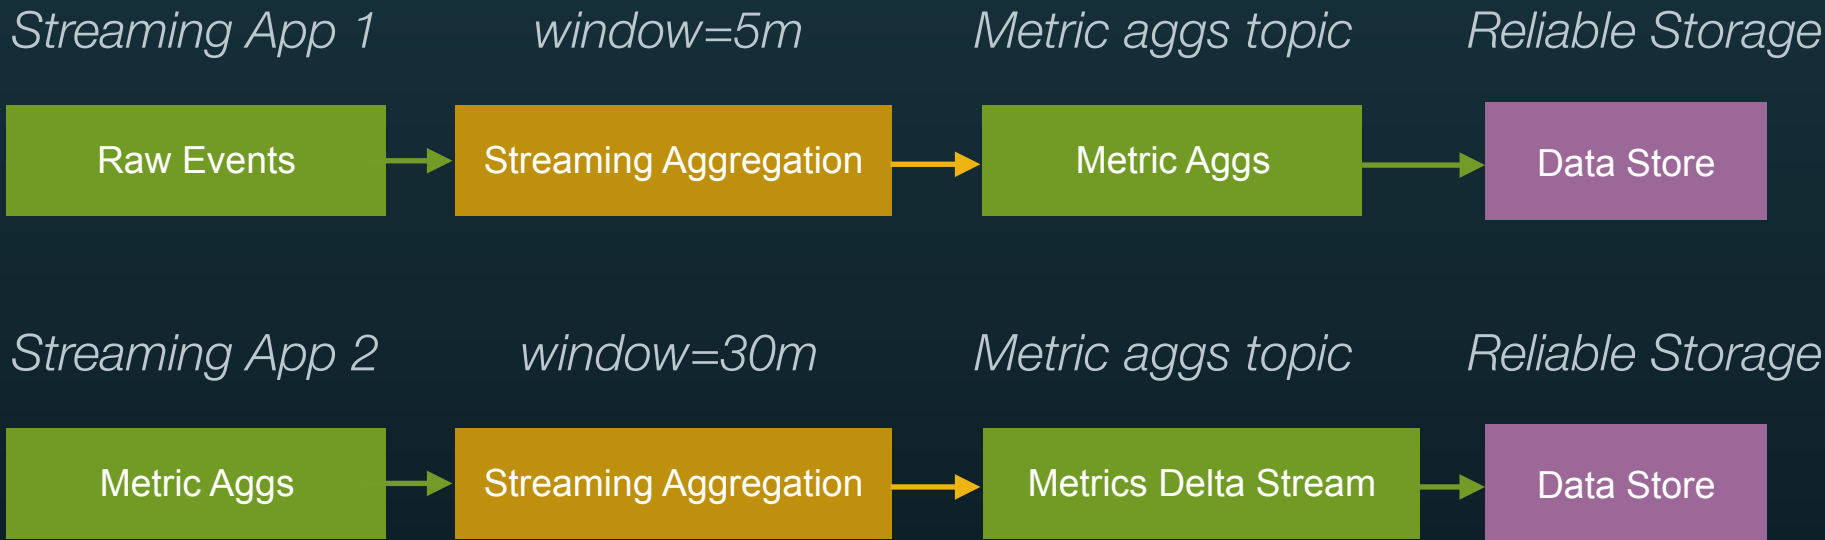
Monadic Systems

1. Distributed
2. Mergeable
3. Accurate

```
abstract class InsightsAggregation extends UserDefinedAggregateFunction {  
  type T<: Product  
  val log: Logger  
  
  def productEncoder: Encoder[T]  
  def productSchema: StructType = productEncoder.schema  
  
  lazy val fieldIndex: Map[String, Int] = productSchema.fieldNames.map {  
    | fieldName => productSchema.fieldIndex(fieldName)  
  }.toMap  
  
  val fieldLength: Int  
  
  override def inputSchema: StructType = InsightsAggregation.inputStringR  
  override def bufferSchema: StructType = productSchema  
  override def dataType: DataType = productSchema  
  override def deterministic: Boolean = true  
  
  override def initialize(buffer: MutableAggregationBuffer): Unit = {  
    (0 to fieldLength).foreach { num =>  
      | buffer(num) = 0L  
    }  
  }  
  
  override def merge(cache: MutableAggregationBuffer, state: Row): Unit = {  
    (0 to fieldLength).foreach { num =>  
      | cache(num) = cache.getLong(num) + state.getLong(num)  
    }  
  }  
}
```

Trend Discovery

Anatomy of the Discovery Engine



Trend Discovery

Code Walk Through

/part3/streaming-trend-discovery

[More Documentation](#)

PART THREE: TREND DISCOVERY

Questions?

#streamingdiscovery | @newfront | @odsc

PART FOUR: TESTING AND OTHER PRODUCTION CONSIDERATIONS

How to test, ship, update and monitor
massive streaming applications

#streamingdiscovery | @newfront | @odsc

Testing and Prod

Core Concepts

1. Testing Spark Apps

1. See ['EventMemoryStreamSpec'](#)

We want to ensure that we can test exactly HOW our aggregations will work.

We should have a repeatable pattern that can be employed on all spark applications.

We should be as confident with our Spark applications as we do traditional Enterprise apps

Testing and Prod

Core Concepts

1. [SparkApplicationListener](#)
2. [StreamingQueryListener](#)

1. *Emits all task, stage, job level details back to the driver.*
2. *Emits progress and statistics for the runtime behavior of a streaming query*

You can use these metrics to Emit to DataDog or other system to be able to track the performance and health of your system


Search events to overlay...

Show 2d The Past 2 Days

Completed and Failed Tasks (by app_id)



Streaming Query Progress



Peak Execution Memory (in MB) per task



Avg. Result Size / task



Kafka Consumer Lag




GC Time




executor tasks jvm.gc.time (avg and max) in seconds




Inserted CS Records



Heap Used (p99) on shared workers



CPU Time per Task



Kafka Consumer (ColSummary)



Spilled Memory Bytes per Task



#streamingdiscovery | @newfront | @odsc

Testing and other production considerations

Code Walk Through

```
/part3/streaming-trend-discovery/.. /listeners/ .. /tests
```

PART FOUR: TESTING AND OTHER PRODUCTION CONSIDERATIONS

Questions?

#streamingdiscovery | @newfront | @odsc